

DATA SCIENCE

CLASS 3: INTRO TO DATA SCIENCE

0. WHAT IS A DATA SCIENTIST?

I. HOW DATA SCIENTISTS ADD VALUE

II. THE DATA SCIENCE WORKFLOW

III. QUALITIES OF A GOOD DATA SCIENTIST

0. WHAT IS A DATA SCIENTIST?

WHAT IS YOUR DEFINITION?



Zvi

@nivertech



Follow

"Data Scientist" is a Data Analyst who lives in California.

↩ Reply ↻ Retweet ★ Favorite ⋮ More

RETWEETS

140

FAVORITES


40



9:55 PM - 14 Mar 2012

WHAT IS A DATA SCIENTIST?

6



Josh Wills

@josh_wills



Follow

Data Scientist (n.): Person who is better at statistics than any software engineer and better at software engineering than any statistician.



Reply



Retweet



Favorite



More

RETWEETS

907

FAVORITES

418

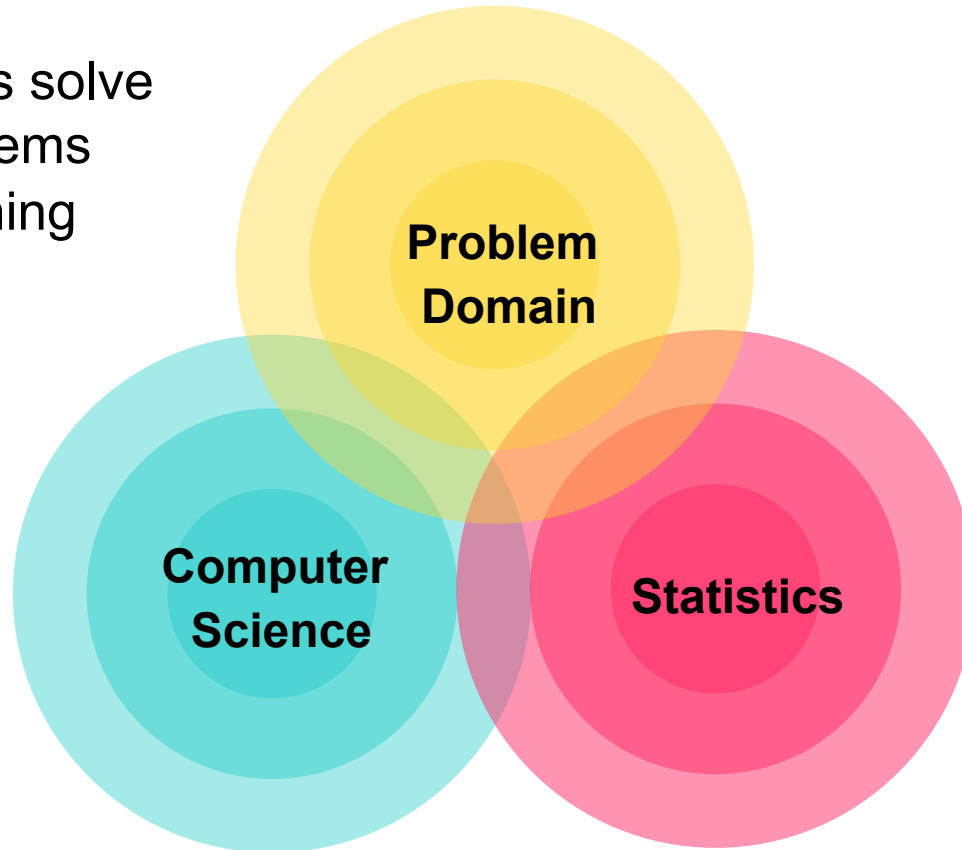


12:55 PM - 3 May 2012

WHAT IS A DATA SCIENTIST?

7

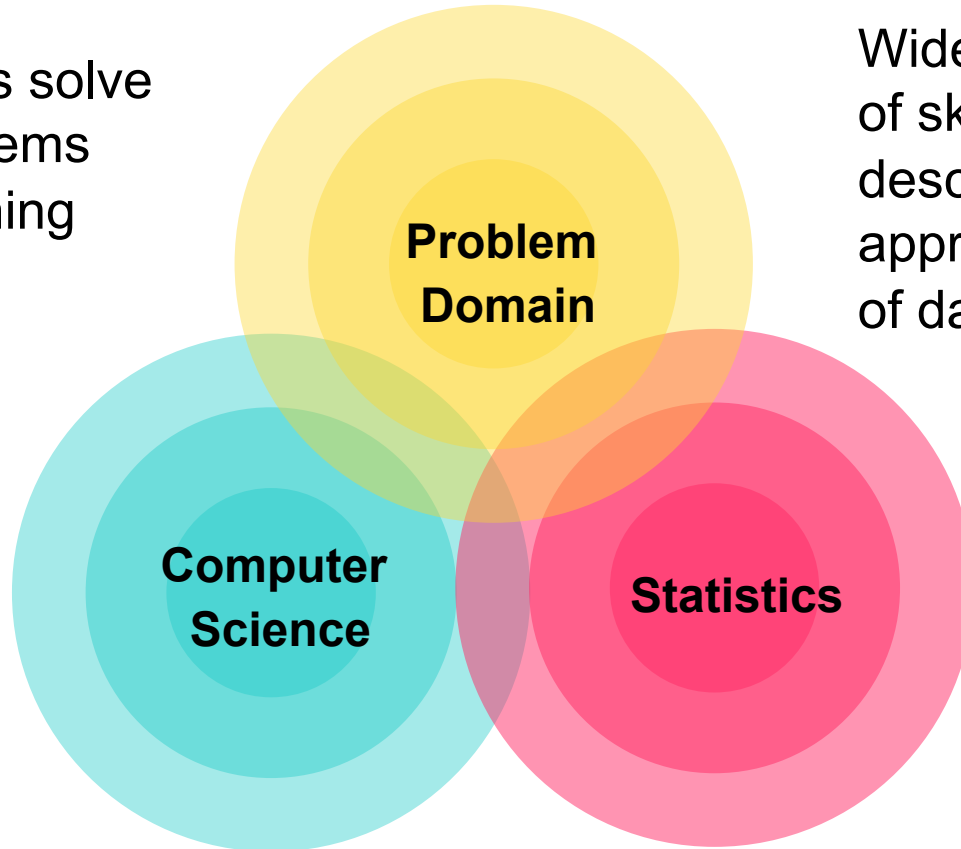
Data Scientists solve complex problems using data mining techniques



WHAT IS A DATA SCIENTIST?

8

Data Scientists solve complex problems using data mining techniques



Wide variance in terms of skillsets: many job descriptions are more appropriate for a team of data scientists

INTRO TO DATA SCIENCE

I. HOW DATA SCIENTISTS ADD VALUE

Data mining techniques generally add value by doing one of four things:

- 1) Predicting the bad
- 2) Identifying the good
- 3) Automating existing processes

Data scientists can be found within many fields: let's look at some additional examples to motivate this course.

EXAMPLE #1: PREDICTING NEONATAL INFECTION

11

Problem: Children born prematurely are at high risk of developing infections, many of which are not detected until after the baby is sick

Goal: Detect subtle patterns in the data that predicts infection before it occurs

Data: 16 vital signs such as heart rate, respiration rate, blood pressure, etc...

Impact: Model is able to predict the onset of infection 24 hours before the traditional symptoms of infection appear



Image: <http://www.babycaretips4u.com/wp-content/uploads/2014/03/premature-baby.jpg>

Case Study: <http://www.amazon.com/Big-Data-Revolution-Transform-Think/dp/0544002695>

EXAMPLE #2: AUTOMATING GOVERNMENT PAPER-PUSHING

12

Problem: Processing disability claims at the Social Security Administration is a time-intensive process, with many claims taking over 2 years to adjudicate

Goal: Automate the approval of a subset of the “simplest” disability claims

Data: Free text in the claims form

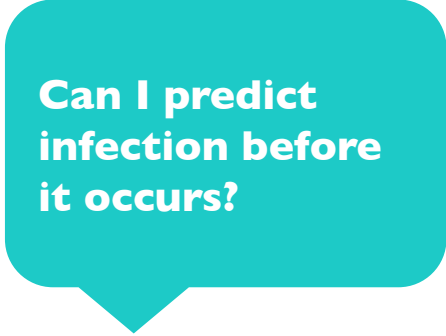
Impact: Able to fully automate 20% of the simplest claims. Rating accuracy of the algorithm is higher than the average claims examiner.



II. THE DATA SCIENCE WORKFLOW

- 0. Define the problem / question**
- I. Identify and collect data**
- II. Explore and prepare data**
- III. Build and evaluate model**
- IV. Communicate results**

0. DEFINE THE PROBLEM / QUESTION



Can I predict infection before it occurs?




Can I predict claim approval from the start of the process?

I. IDENTIFY AND COLLECT DATA

**Vital Areas:
Heart Rate,
Blood Pressure,
etc...**

**Want to collect
all data on the
claim form
(mostly free
text)**

II. EXPLORE AND PREPARE DATA



**Aggregate data
at the minute
level**



**Cluster like
words**

III. BUILD AND EVALUATE MODELS

**Compare
Decision Tree
with Logistic
Regression**

**Start with Naïve
Bayes Classifier**

IV. COMMUNICATE RESULTS



**Create custom
dashboard for
doctors and
nurses**




**Create report
and dashboard
proof of concept**

III. QUALITIES OF A GOOD DATA SCIENTIST

**ASKS
RATIONAL
QUESTIONS**

**UNDERSTANDS
THE PROS & CONS
OF DIFFERENT TECHNIQUES**

**STATISTICIANS, LIKE
ARTISTS, HAVE THE BAD
HABIT OF FALLING IN LOVE
WITH THEIR MODELS
– GEORGE BOX**



**COMMUNICATES
CLEARLY**

**RETAINS
INTELLECTUAL
HUMILITY**

